

What you actually pay for

From the unit on the price page to the number that decides whether it was worth running.

COMPONENT

ARTIFACT

COST

TOKEN BILLING

charged by consumption, not by seat — so the bill scales with how much the thing is used

01

Token

the unit everything is counted in

Everything downstream is denominated in these: the price, the window, the rate limits. Not words, and not characters.

A WORD FRAGMENT

02

Token pricing

three rates, not one

Input and output are charged at different rates, and cached input cheaper again — which is where most of the saving available to you actually is.

OUTPUT COSTS THE MOST

03

The levers

what changes the rate you pay

Prompt caching

A repeated prefix processed once rather than on every call.

Batch inference

Roughly half price for giving up any promise about when.

Provisioned throughput

Guaranteed capacity by the hour: a variable cost made fixed, in both directions.

04

Cost per task

what one finished job cost

Retries, tool calls and the runs that failed, counted against the jobs that completed. An agent making forty calls to finish once is priced by the finish, not by the call.

NOT THE PRICE OF A CALL

WHAT YOU ARE SOLD

Subscription tier — The packaging that decides which capabilities you get at all — and the usual reason two people describe the same product differently.

WHERE IT STOPS

Usage limit — The ceiling a plan puts on consumption over a period. A rate limit shapes traffic; this one stops it.

NOT PAYING PER TOKEN AT ALL

Self-hosting

Run it on infrastructure you control — a hardware decision before it is a software one.

On-device inference

The user's own phone or laptop, so nothing is sent anywhere.

Per-token prices compare models. Cost per task compares decisions — and the two can point in opposite directions, because the cheaper model is often the one that needs more attempts.