

How you know it got better

The four claims you will be shown, and the question each one is hiding.

WHAT YOU WILL BE TOLD

It scores 87 on the benchmark

A benchmark is shared, which is the point of it — and being shared makes it a target that models get optimised against rather than measured by.

Leaderboard

It is top of the leaderboard

The most-read artifact in the field, and one that stops measuring the thing it ranks once enough people are optimising for the rank.

A model scored the output as better

Cheap enough to run on everything, which is why it is everywhere — and carrying biases that the paper naming the technique documented in detail.

Autorater evaluation

The fix works — we tried it

Trying it once is not a result when the same input does not reliably produce the same output. Nothing here regresses loudly; it regresses on the case nobody re-ran.

WHAT TO ASK ABOUT IT

Was the test set in the training data?

Test data leaking into training inflates the score without improving anything, and it is now the normal condition rather than an accident. Ask when the set was published relative to the cutoff.

Does it win on your task?

A public ranking answers a question somebody else asked. The only test that settles your decision is one built from your own inputs and run repeatably.

Regression testing

Was the judge checked against people?

A scorer is itself a model with an error rate. Ask what it was calibrated on, and whether anyone has measured how often it agrees with a human on this kind of output.

Citation precision and recall

Can you see what it actually did?

Record the run step by step. Without a trace, “why did it do that” has no answer, and a fix cannot be shown to have addressed the cause rather than the symptom.

Red teaming

THE WORD DOING THE WORK

Evaluation — A repeatable test. Repeatable is the whole of it: a number nobody else can reproduce is an anecdote with a decimal point.

WHAT A SCORE CANNOT TELL YOU

Factuality — Whether the output is true, as distinct from whether it sounds right — a property to aim at, not a number to report.

FINDING FAILURES ON PURPOSE

Red teaming — Deliberately trying to make it misbehave, rather than checking that it behaves. The two produce very different reports.

Every number here was produced by someone who chose the test. The question is never “what did it score” but “on what, against whom, and could I run it again”.